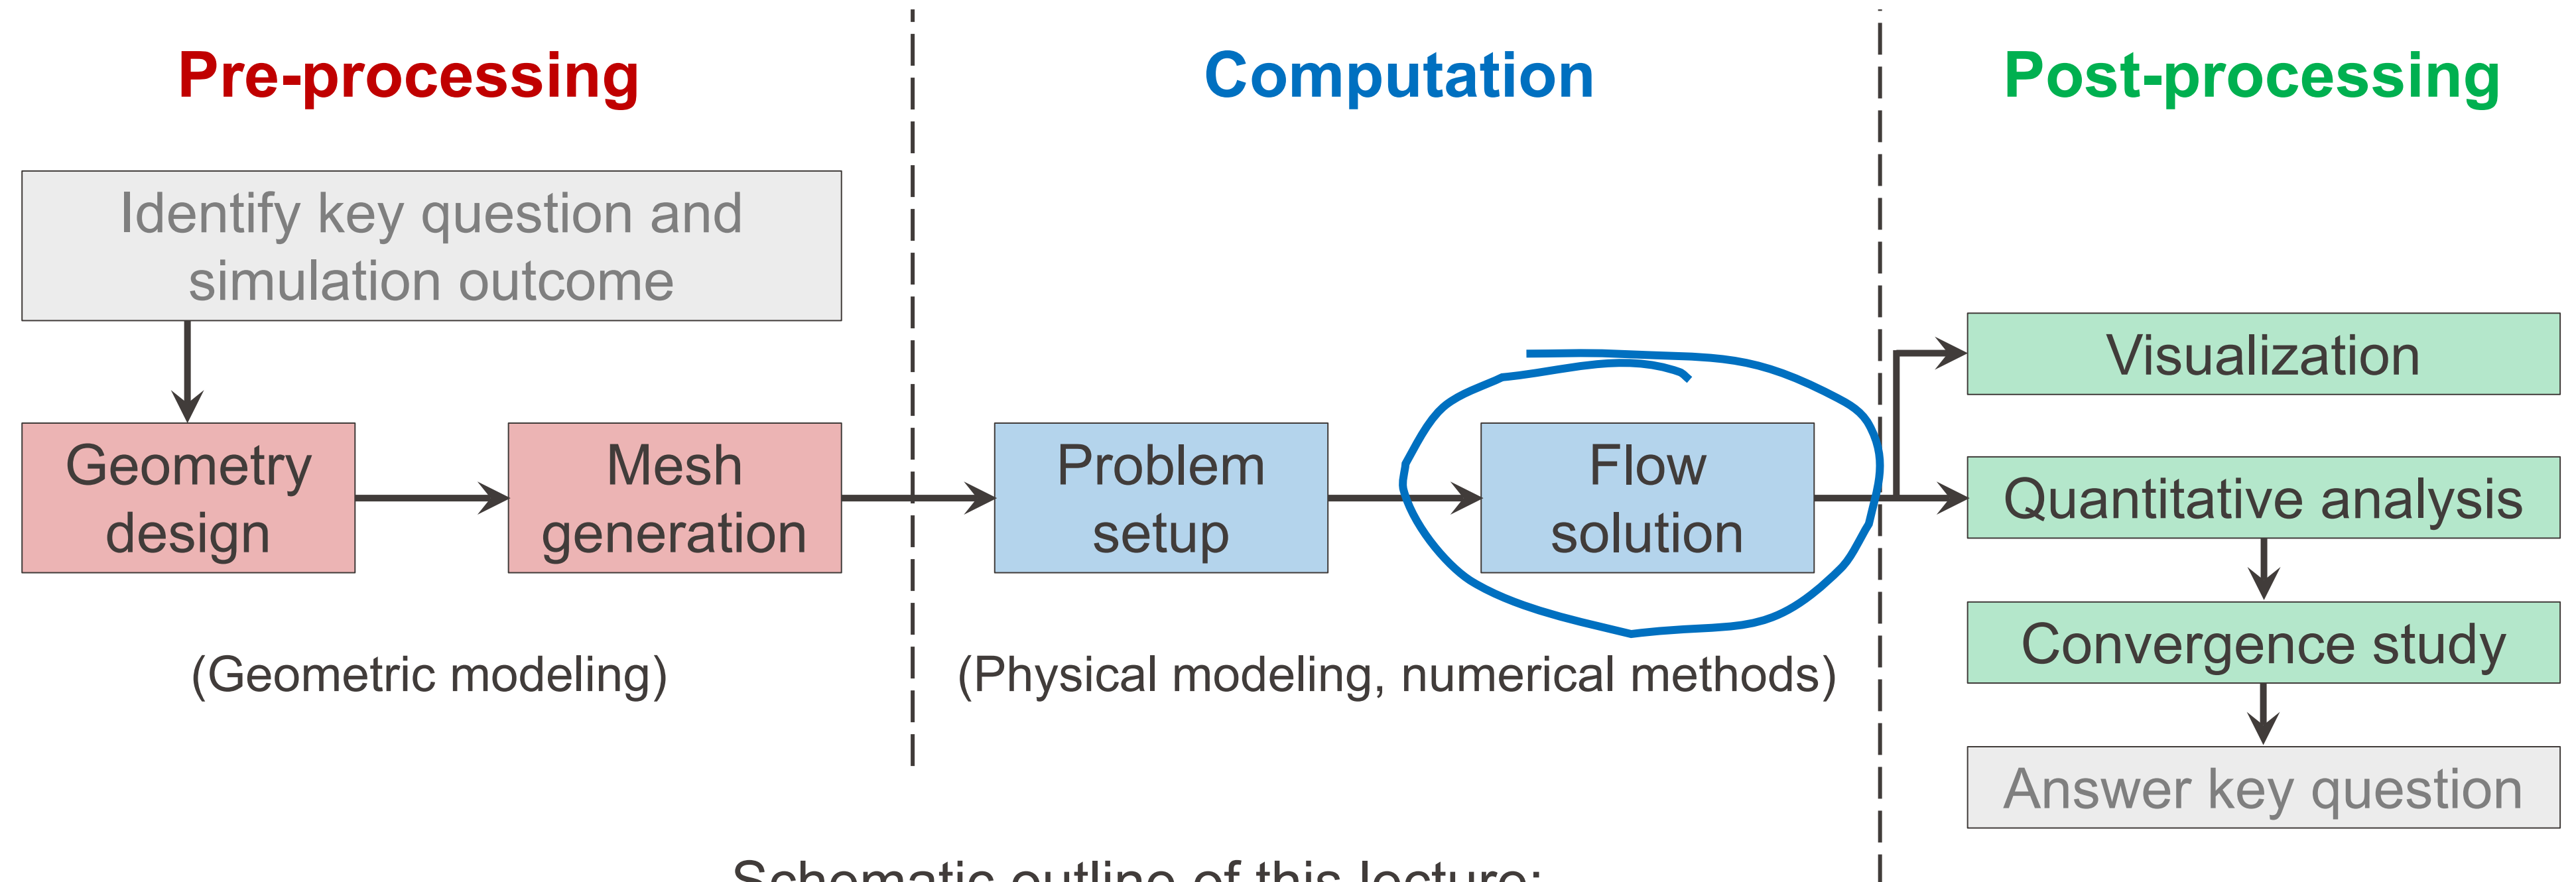


Parallel computation

Numerical Flow Simulation

Numerical simulation workflow




Schematic outline of this lecture:

1. Why parallel CFD?
2. Implementation in Fluent
3. Mesh partitioning
4. Parallel performance

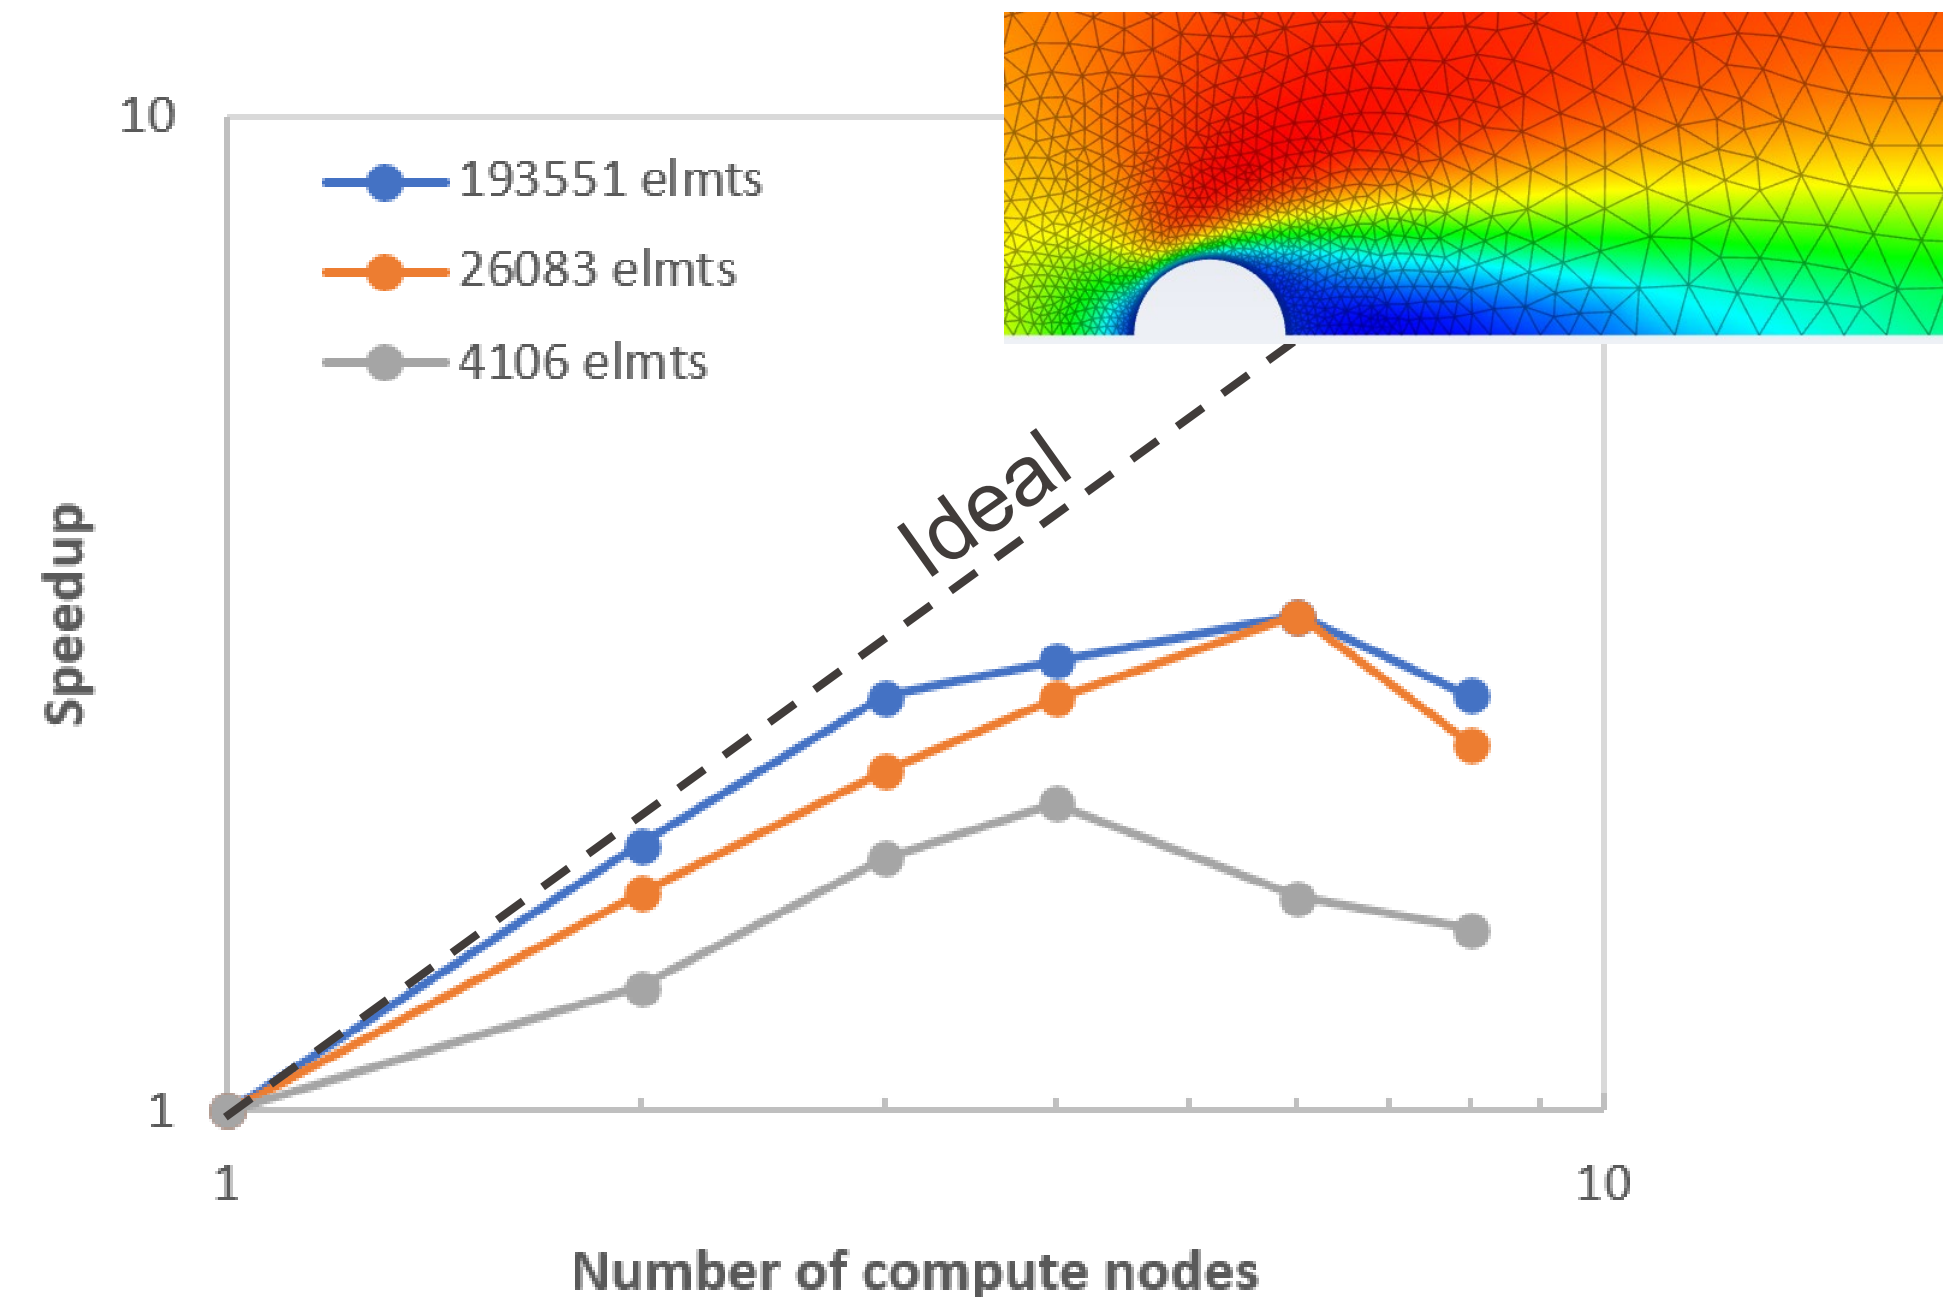

Why parallel CFD?

- Basic idea: distribute the calculations among several processors that work simultaneously → **more memory available** & **shorter calculation time**.
- Perform **larger** simulations
 - Larger meshes
 - More complex physics
- Perform simulations **faster**
 - Many designs (virtual prototyping).
 - Shorter turn-around time (e.g. from days to hours, for several-million cell meshes)
- Take advantage of state-of-the-art high-performance computers
 - Fast computation
 - Large memory



Why parallel CFD?

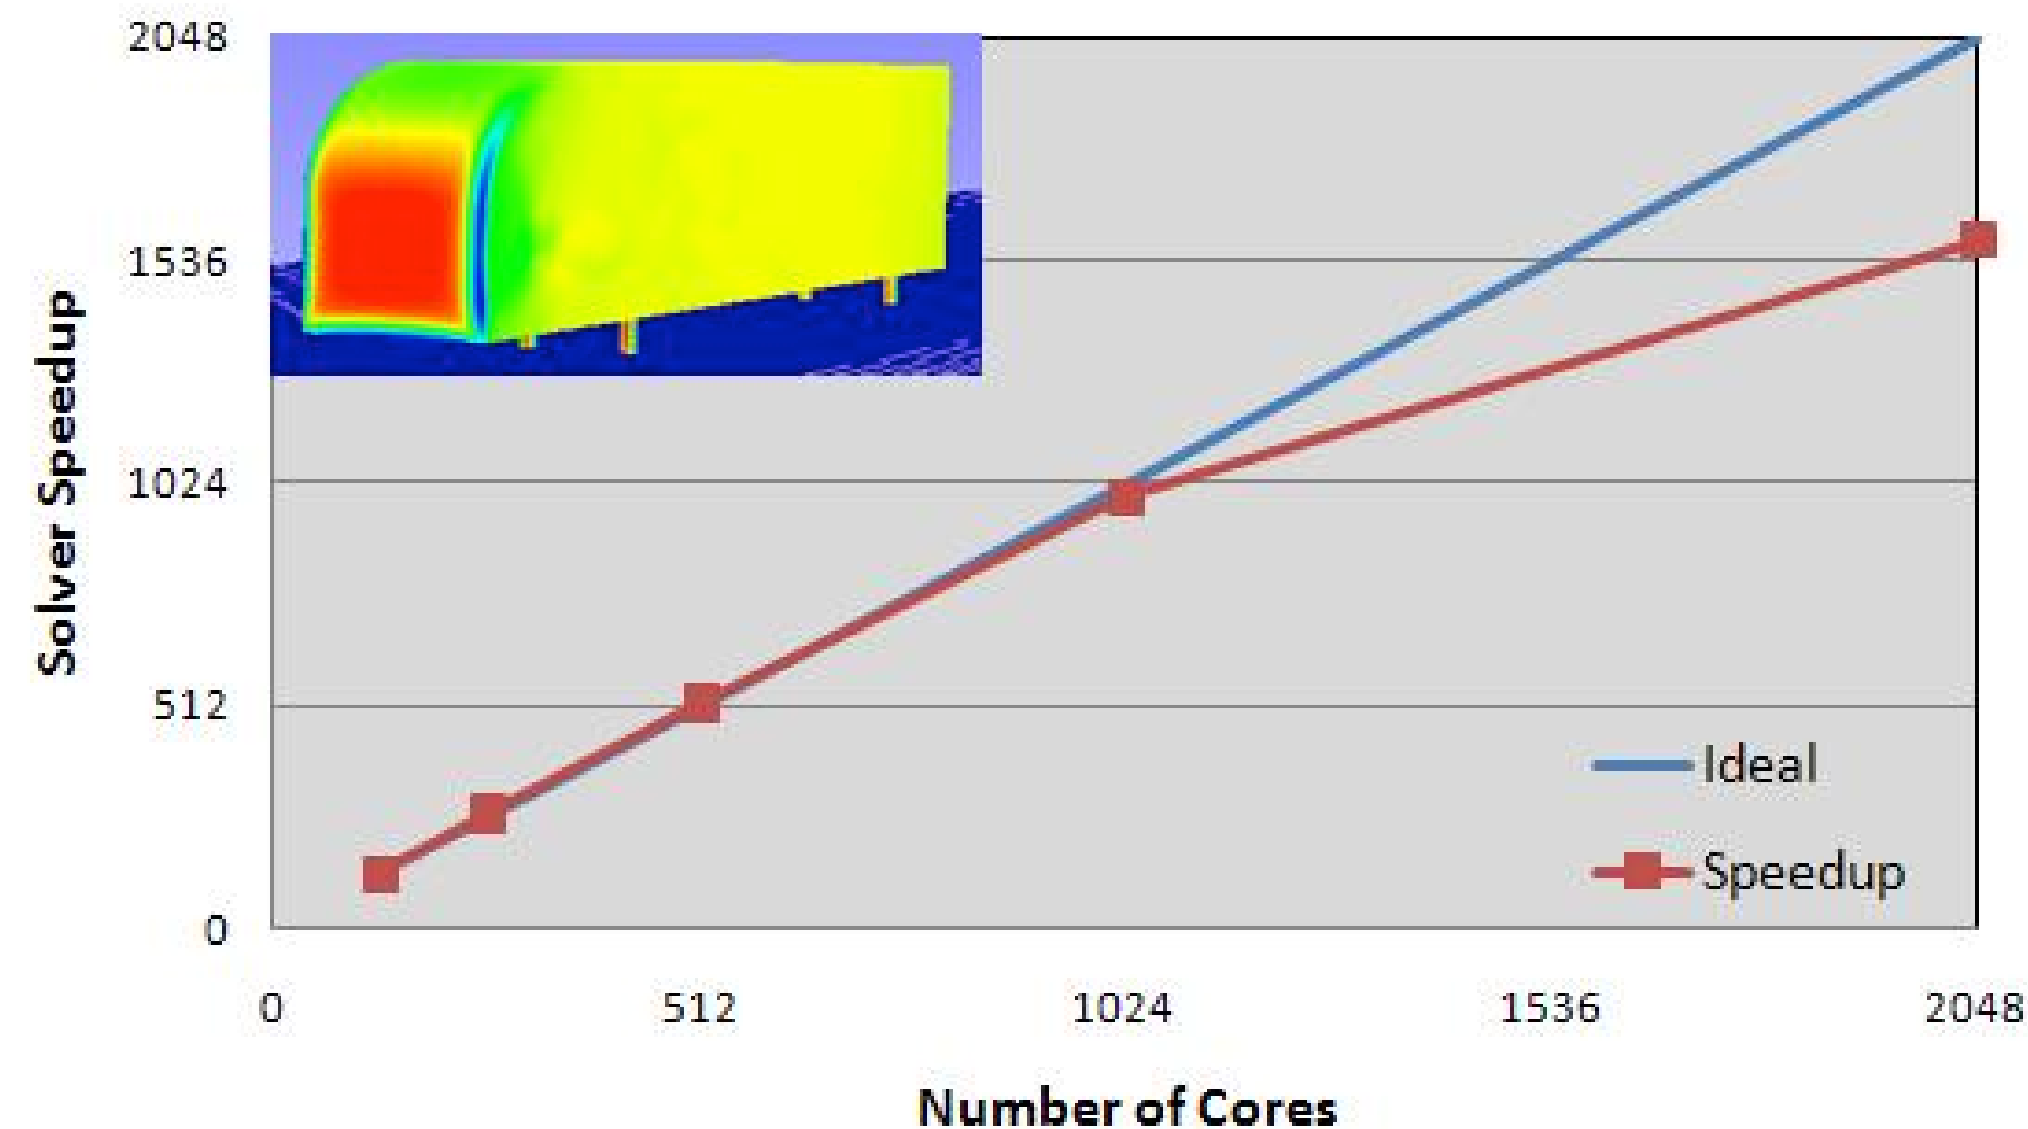
- When should parallel computation NOT be used?
 - When the problem is too small: too little work per processor and too much communication between processors → inefficient (or even slower!).



Why parallel CFD?

- When should parallel computation NOT be used?
 - When the problem is too small: too little work per processor and too much communication between processors → inefficient (or even slower!).

Truck Benchmark (111 Million Cells)

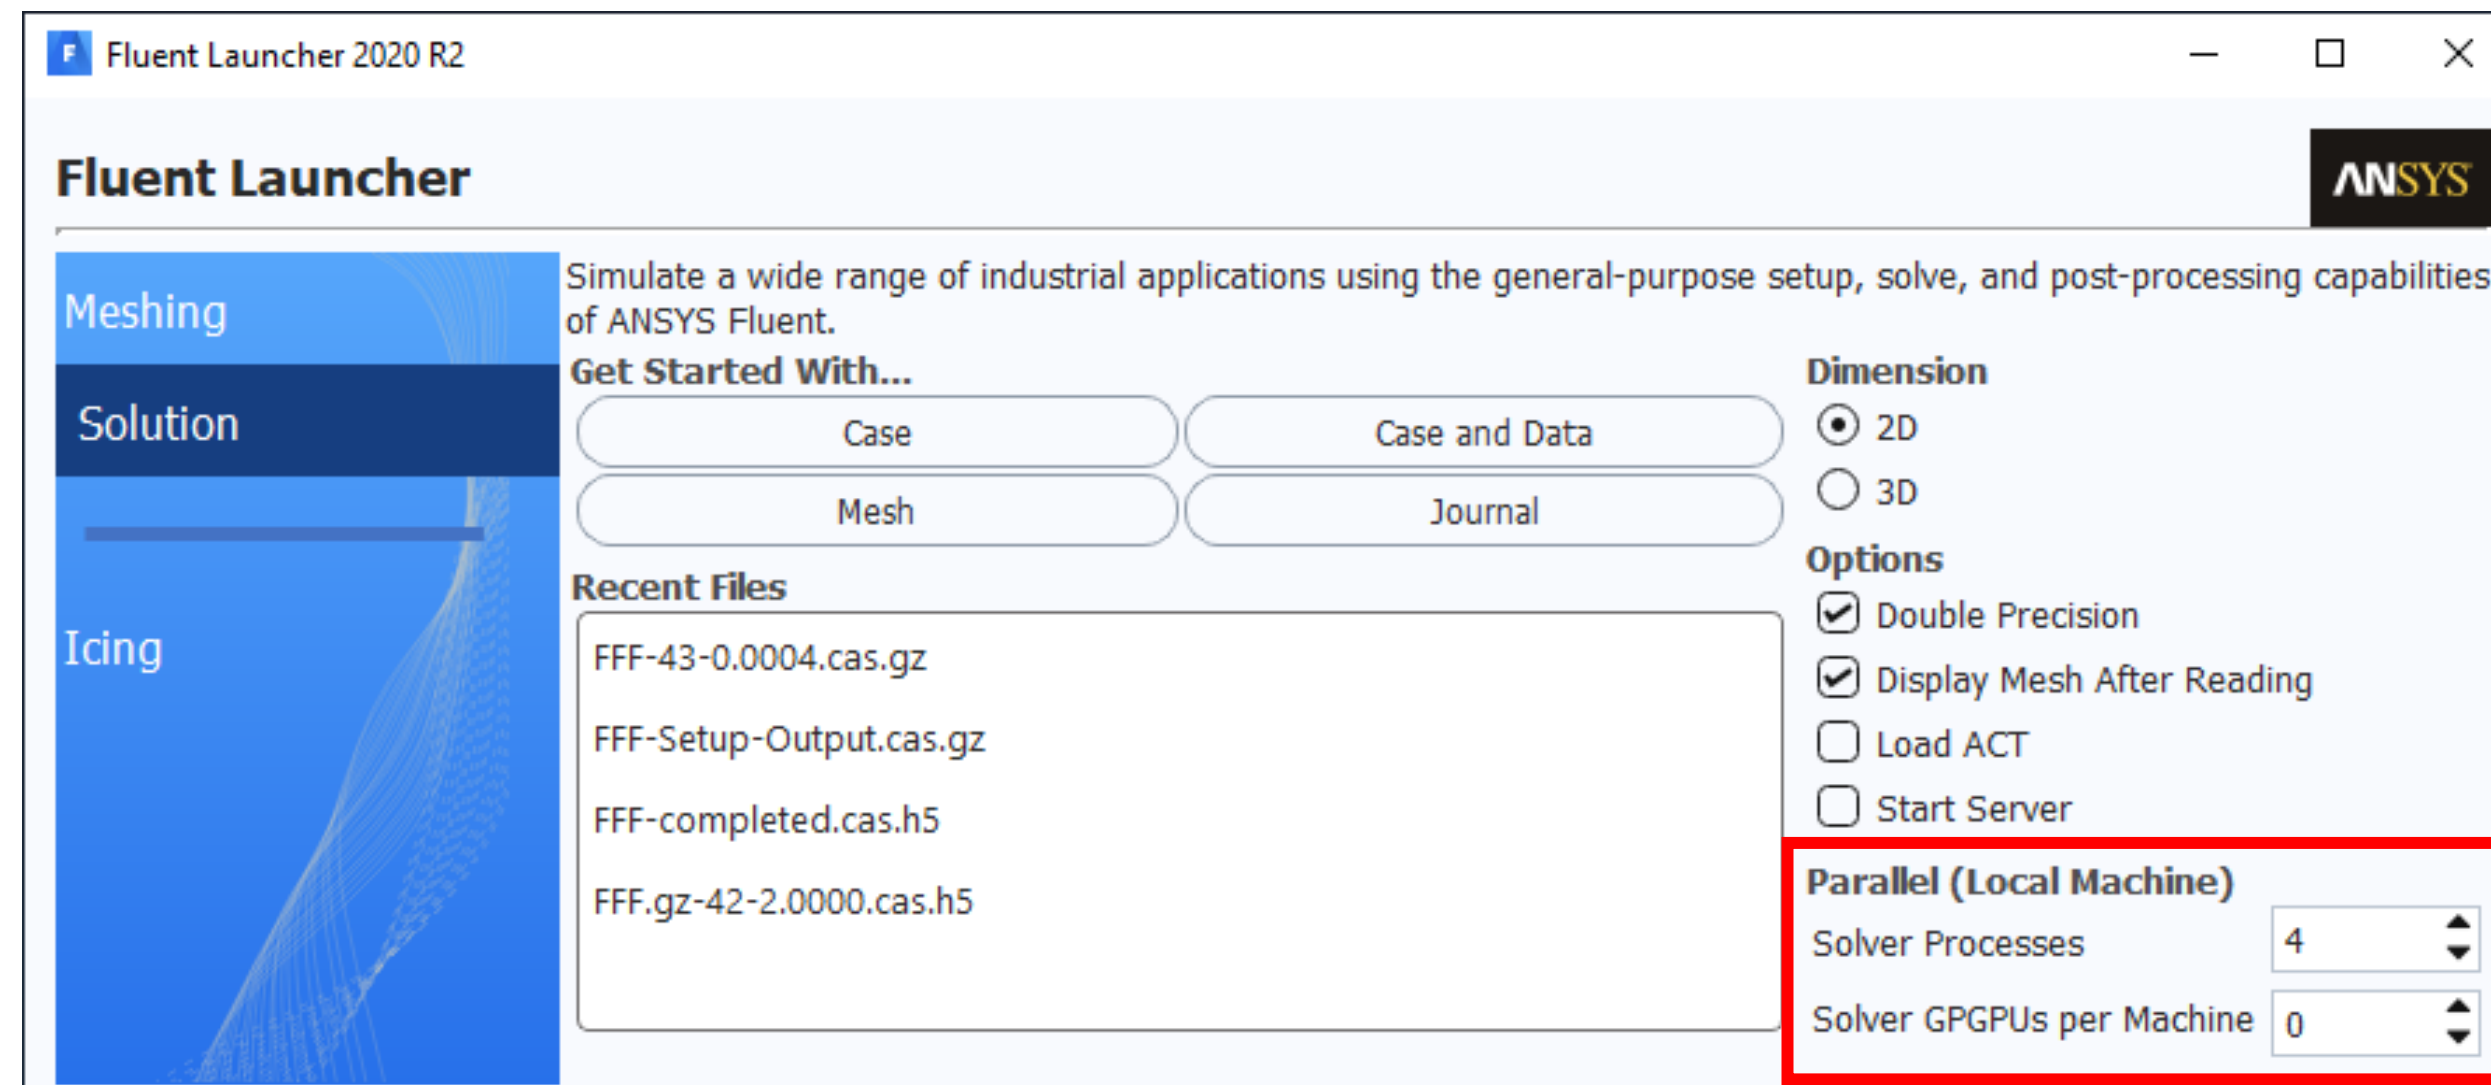


Why parallel CFD?

- When should parallel computation NOT be used?
 - When the problem is too small: too little work per processor and too much communication between processors → inefficient (or even slower!).
 - When the computer resources are not appropriate (e.g. insufficient communication bandwidth) → must check parallel performance (see later).
 - When the user doesn't know what they are doing...

Implementation in Fluent

- Assumption:
 - Users of commercial software are interested in shortening turn-around time, not in parallel computing itself.
 - May want to use software on different serial/parallel computer systems.
- Implementation:
 - As “transparent” as possible: detailed knowledge of parallel computing not required.
 - The code should be highly portable.



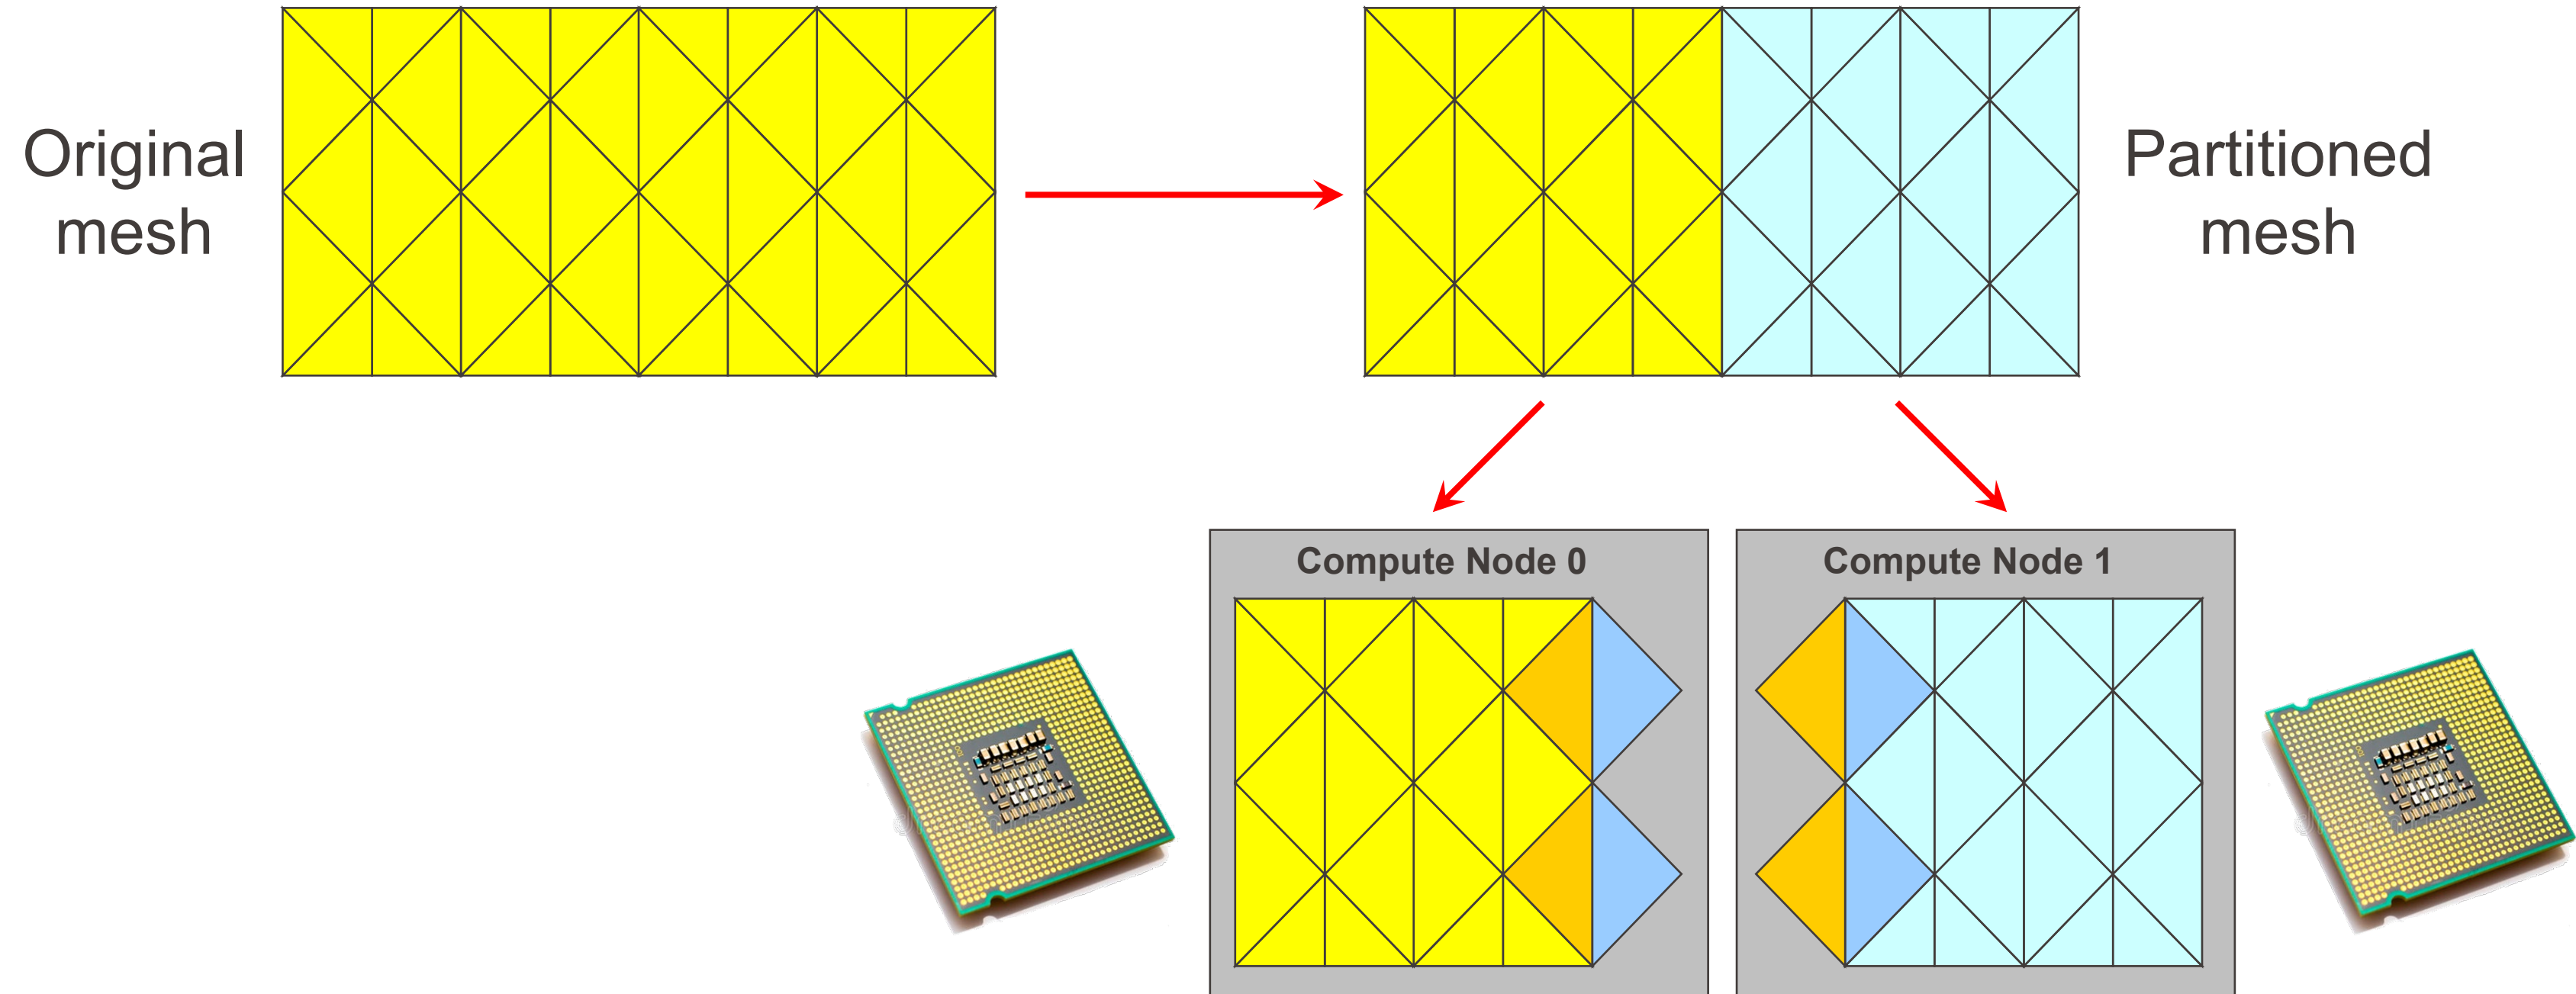
Basically, the user only has to choose the number of processors.

Implementation in Fluent

- Memory model: supports both shared-memory and distributed-memory parallel computers (see later).
- Based on domain decomposition: computational mesh partitioned into sub-meshes, each partition assigned to a different compute node (see later).
- Communication between processors: message passing
 - Scalable
 - Complete control over communication for better performance optimization
 - MPI standard has emerged with wide support
- Client/server model: separation of front-end tasks (GUI, visualization) and solver.

Mesh partitioning


- Basic concept: computational mesh automatically partitioned into several sub-meshes (partitions), each partition assigned to a different processor/node that solves the equations only in its own sub-mesh.



Distributed sub-meshes (with “ghost cells”
to exchange information about neighboring cells)


Mesh partitioning

- Wide range of algorithms:
 - Coordinate bisection,
 - Recursive spectral bisection,
 - Multi-level graph partitioning...
- Goals for optimal algorithm:
 - Partitions with similar number of cells (similar load on each processor)
 - Simply-connected partitions (no holes)
 - Minimize data transfer (interface boundaries with short length and small number of faces)




Mesh partitioning


- Fluent can use the following algorithms:
 - Cartesian axes; Cartesian x,y,z ; Cartesian strip
 - Principal axes (default); principal x,y,z ; principal strip
 - Polar axes; polar r,θ (only in 2D)
 - Spherical axes; spherical ρ,θ,φ (only in 3D)
 - METIS (multi-level graph partitioning)



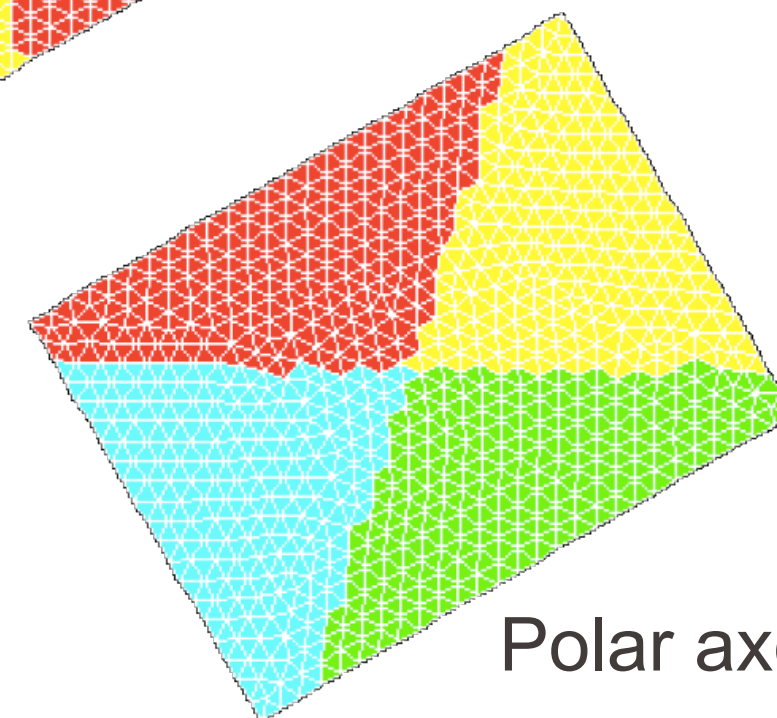
Cartesian axes




Cartesian x



Principal strip

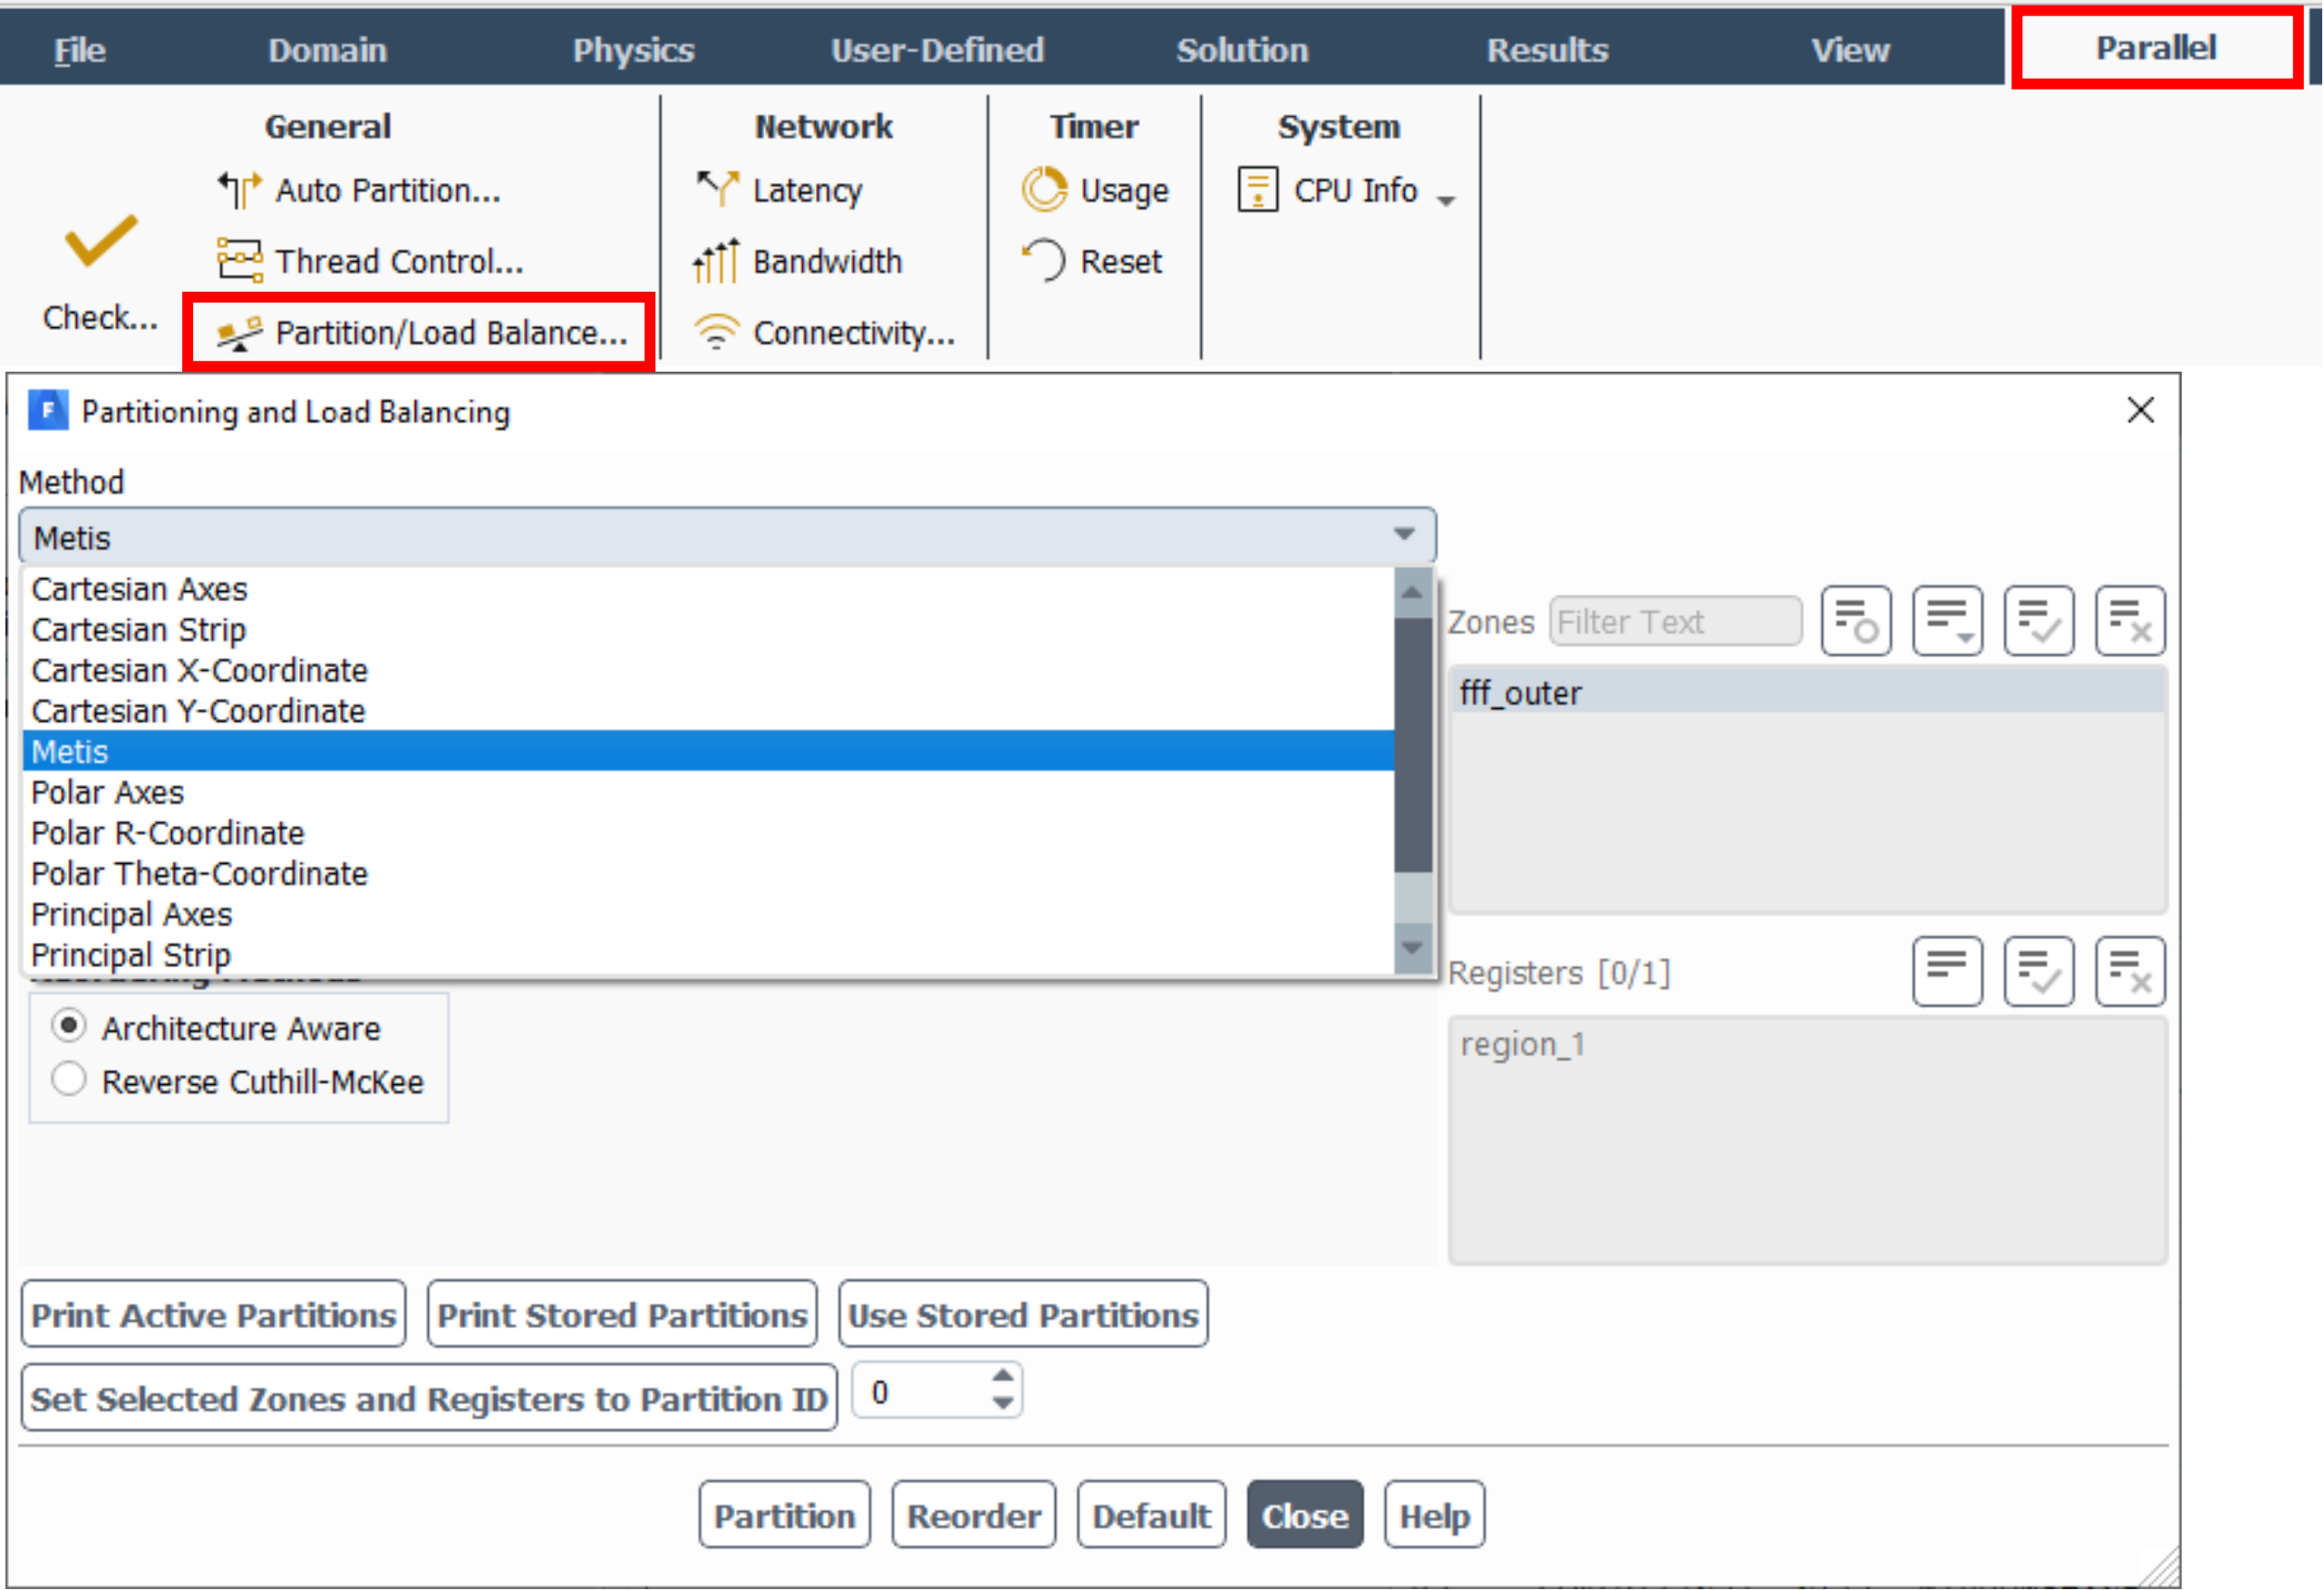


Polar axes



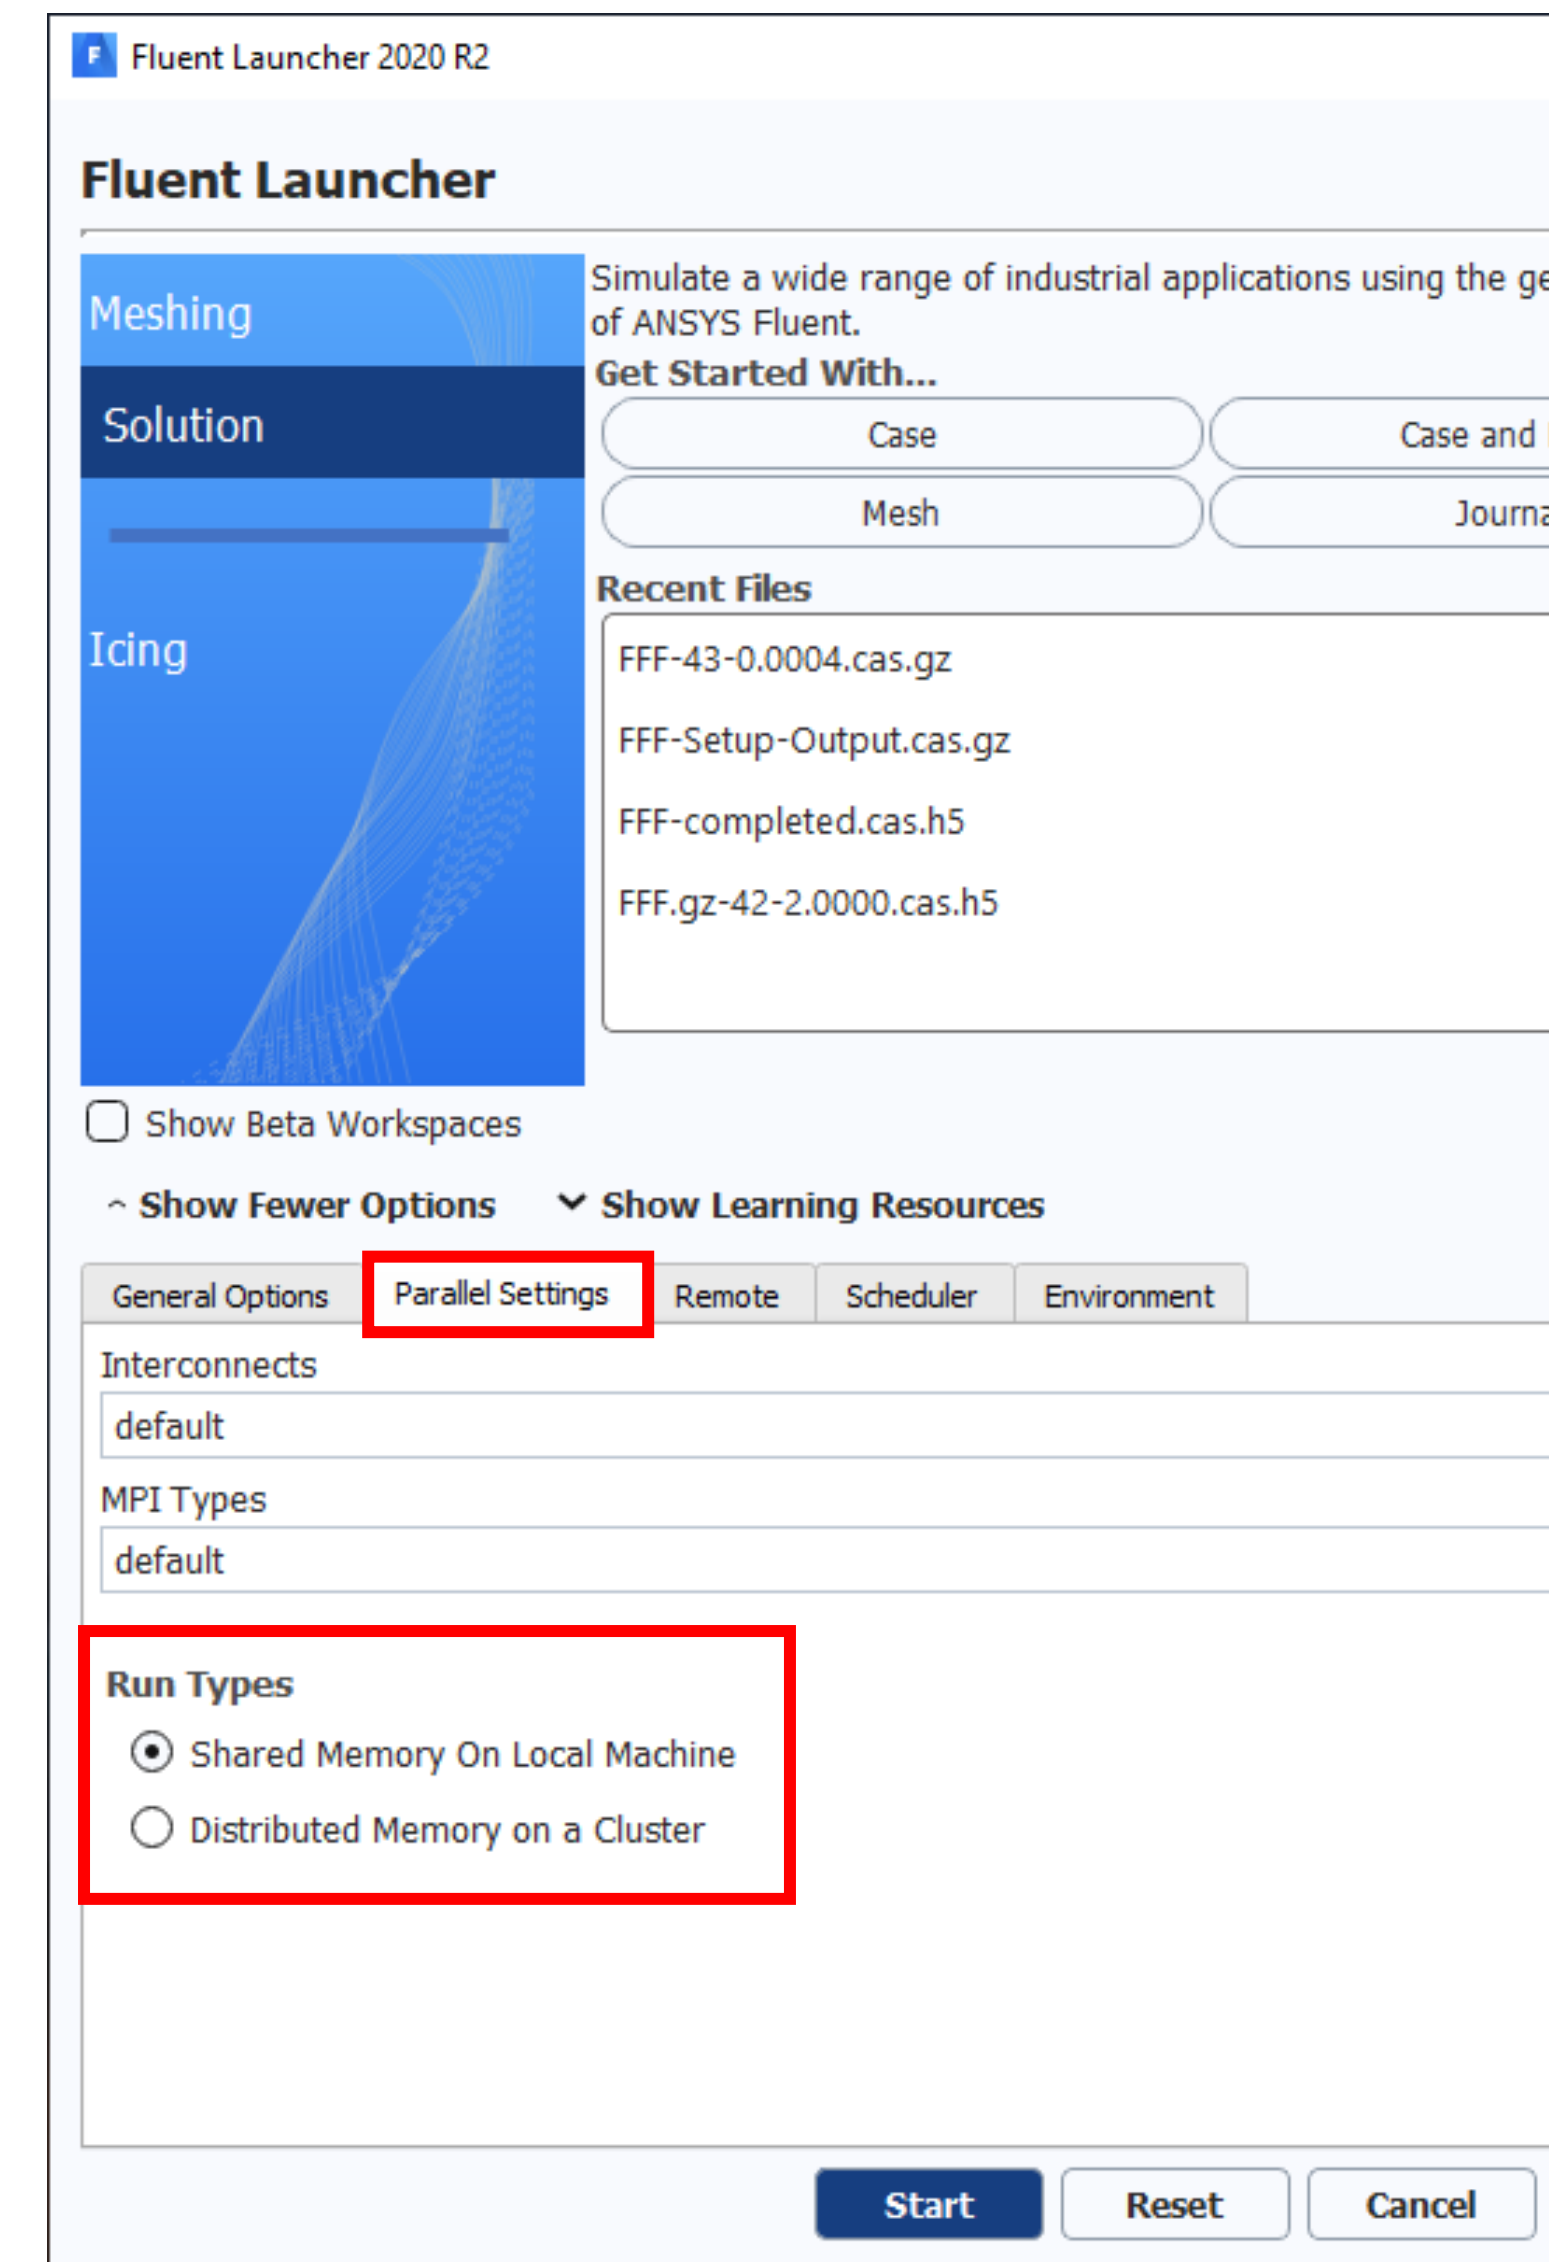
Principal axes

Mesh partitioning




Architecture

- Fluent can run parallel calculations on 2 types of architecture:
 1. **Shared memory: one machine** with several processors (from simple PC/laptop to high-performance workstations)
 2. **Distributed memory: cluster** of several machines (e.g. EPFL clusters)



EPFL SCITAS clusters

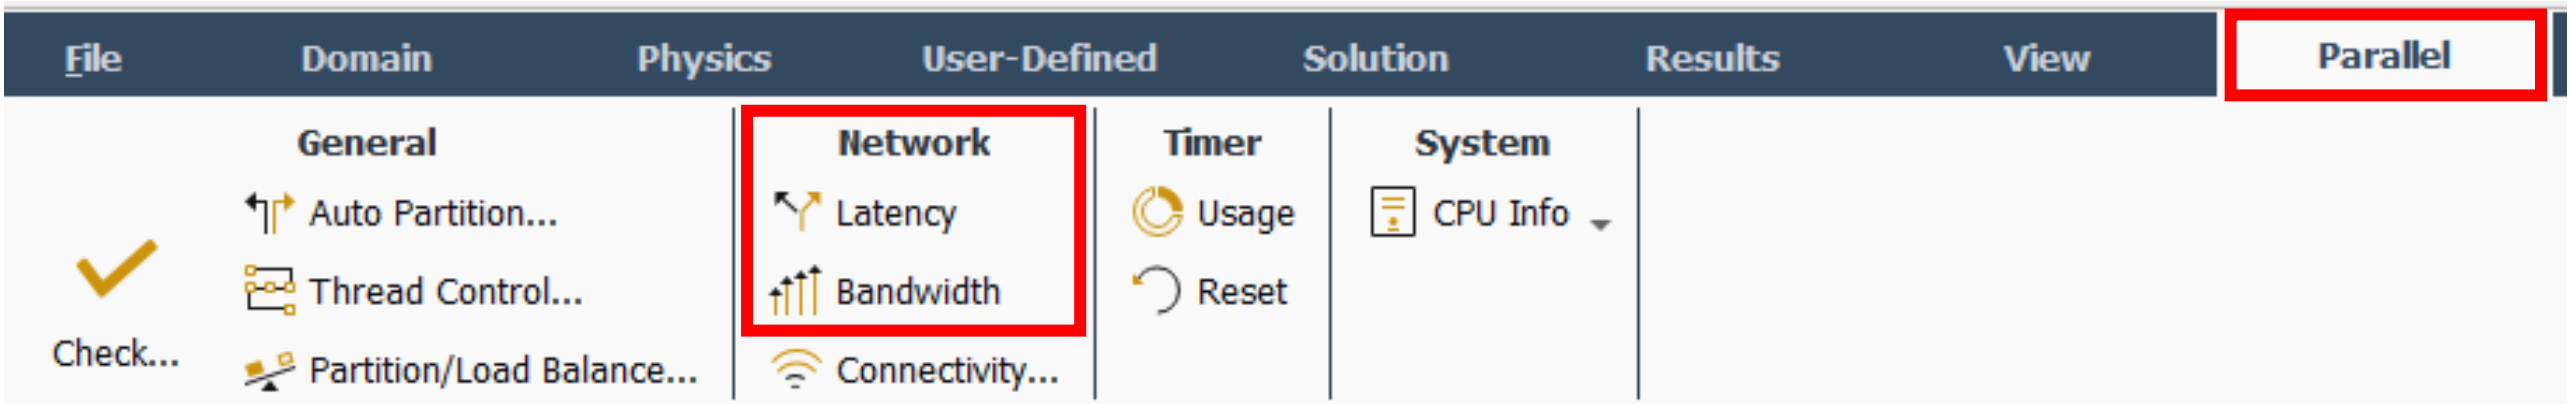
- Several high-performance clusters are available at EPFL.
- Some of them may be available for student projects.



<https://www.epfl.ch/research/facilities/scitas/>

Parallel performance

- Two important characteristics of parallel communication:
 - Bandwidth:** transmission speed = max. amount of data sent in given time [MB/s]
 - Latency:** time to send data = delay [ms] or [μ s]
- Fluent can measure bandwidth and latency between processors.



Bandwidth (MB/s) with 5 messages of size 4MB

ID	n0	n1	n2	n3	n4	n5
n0		111.8	*55	111.8	97.5	101.3
n1	111.8		69.2	98.7	111.7	*51
n2	54.7	69.2		72.9	104.8	*45
n3	111.8	98.7	72.9		64.0	*45
n4	97.6	111.7	104.8	*64		76.9
n5	101.2	50.9	45.5	*45	76.9	

Latency (usec) with 1000 samples

ID	n0	n1	n2	n3	n4	n5
n0		48.0	48.2	48.2	48.3	*50
n1	48.0		48.2	48.3	48.3	*48
n2	48.2	48.2		48.8	49.1	*53
n3	48.2	48.3	*49		48.6	48.5
n4	48.3	48.3	49.1	48.6		*50
n5	49.7	48.5	*53	48.5	49.7	

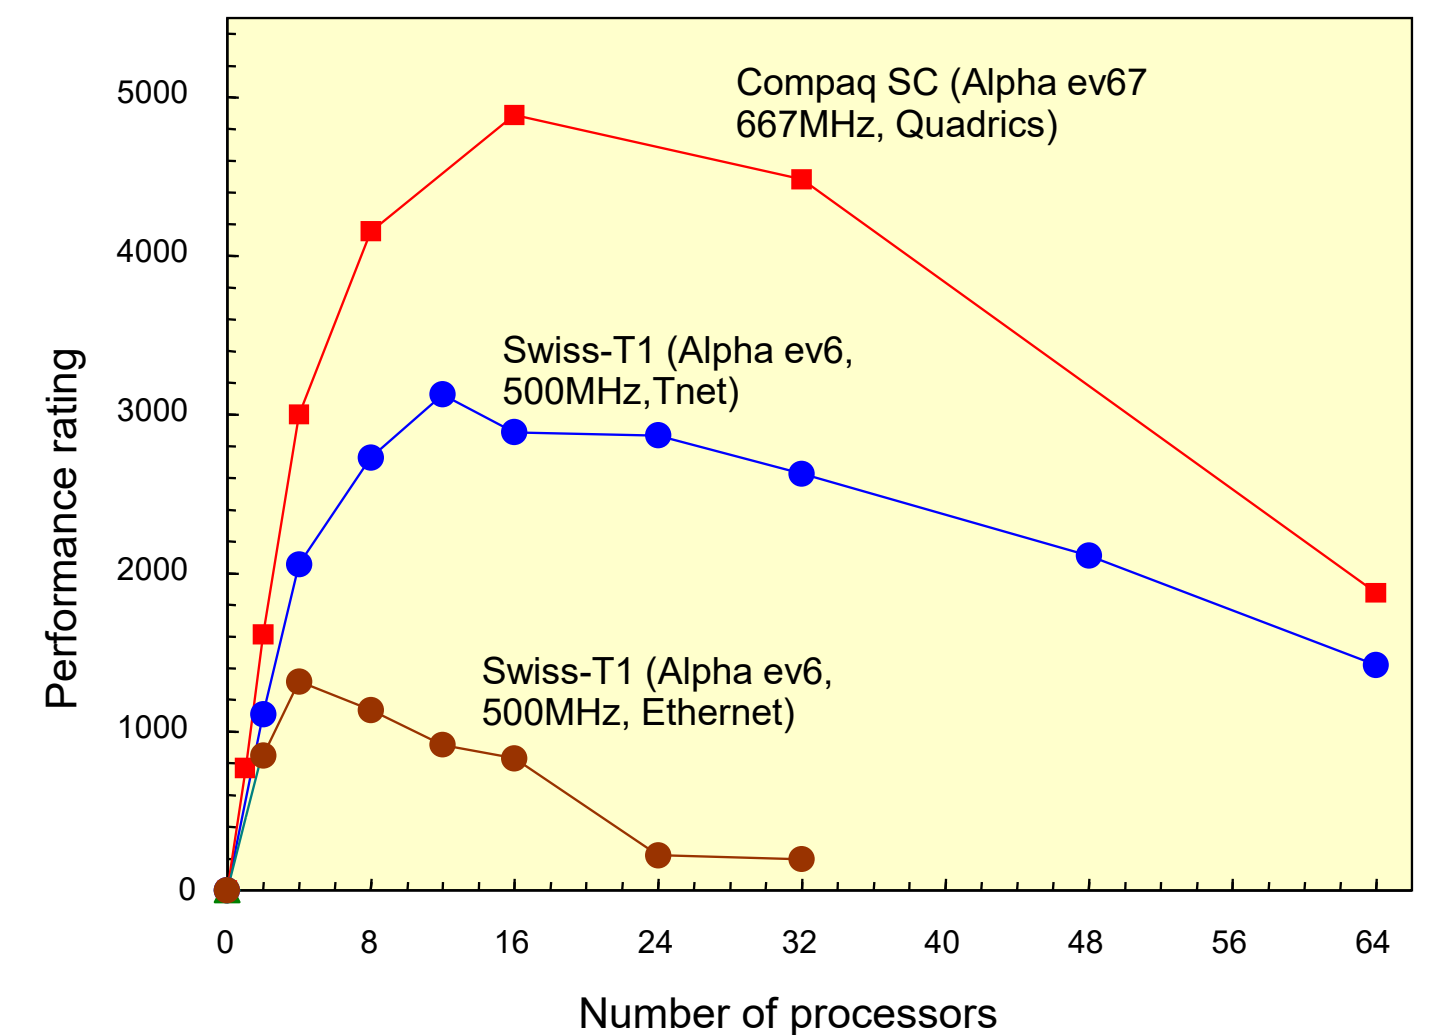
Parallel performance

- Both bandwidth and latency influence parallel performance.
- Latency tends to have a greater effect (over a range of common problem sizes and available networks)
 - Scalability of Fluent largely determined by scalability of its linear equation solver. Multigrid limited by latency because requires several (small) message exchanges.
 - A low-latency interconnect can greatly improve performance scalability of Fluent.
- Performance also depends on:
 - Problem size
 - Physical modeling
 - Numerical method

Parallel performance

- Optimal use
 - For a given problem (size, physical modeling, numerical method):
 - Performance scales well for small number of processors
 - Performance will saturate (or decrease) for large number of processors
- Optimal value of number of processors:
 - Increases with problem size
 - Increases (generally) with problem complexity
- Choose wisely:
 - Parallel computer system (bandwidth & latency)
 - Number of processors

3D turbulent flow in a bend, 32'000 hex cells

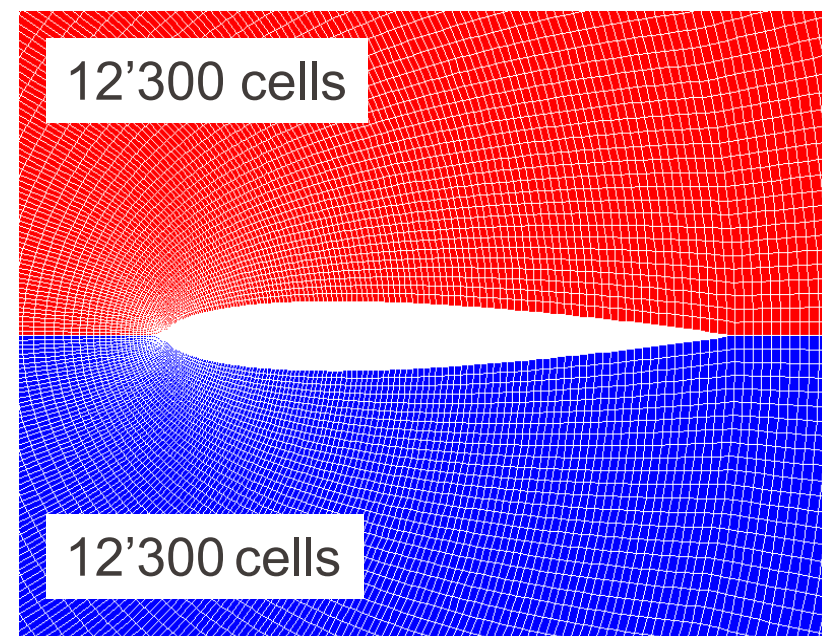


Parallel performance

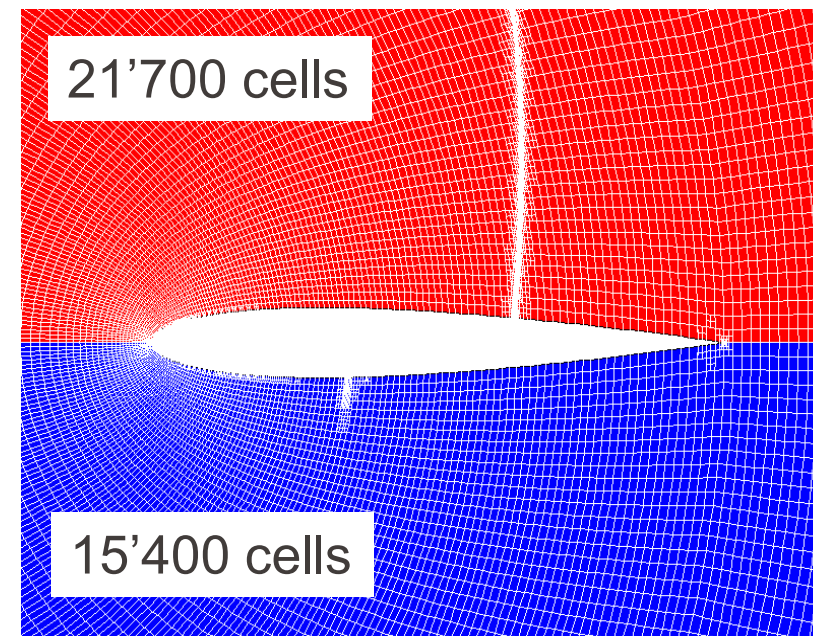
- Dynamic load balancing
 - Automatically redistributes the computational load more evenly among processors (cell migration from one processor's memory to another).
- Load imbalance can result from:
 - Change in performance of some processors due to other processes
 - Non-uniform network performance
 - Changes in mesh distribution from refinement/coarsening
- Associated with performance penalty (disabled by default in Fluent, must be explicitly enabled if desired).

Parallel performance

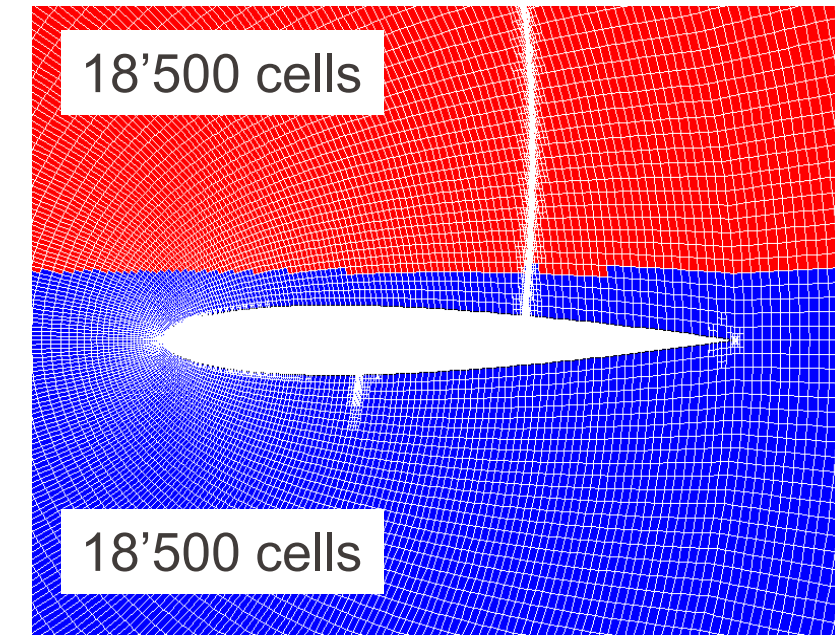
- Dynamic load balancing: example (transonic flow over a NACA 0012 profile)



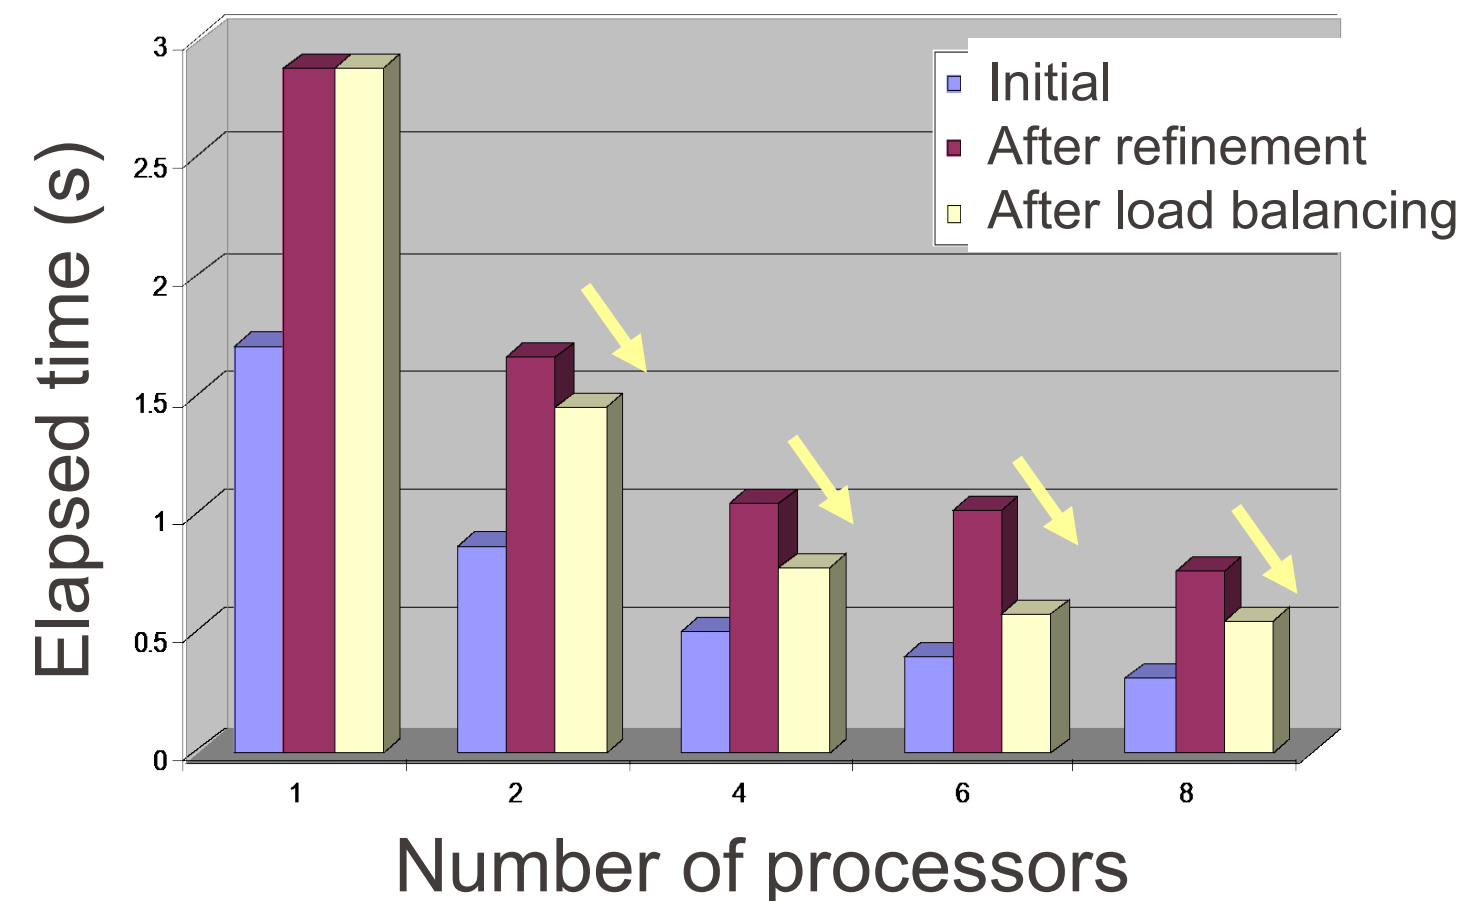
Initial mesh
(here on 2 partitions)



Refined mesh
→ imbalance



Refined mesh
after load balancing



Summary

- Parallel computation allows “large” problems to be solved:
 - Very fine mesh
 - Long-time behavior
 - Complex physics
- Network is characterized by bandwidth (transmission rate) & latency (delay)
- Performance of parallel computation depends on:
 - Numerical problem
 - Network properties
 - Compute node properties
- To learn more: course “Parallel and high-performance computing” (MATH-454)
- Remember: numerical simulation has a cost and an energy footprint
→ only run simulations that are needed!